

Matrix x Matrix multiplication

As a step by step instruction has been presented in tutorial 2, here is a time for a stand-alone practice.

Accelerate the serial, element-wise square matrix addition code using cuda kernel.

```
In [1]: %%file matrix_add.cu

// This program computes a simple version of matrix multiplication
// By: Nick from CoffeeBeforeArch

#include <algorithm>
#include <cassert>
#include <cstdlib>
#include <functional>
#include <iostream>
#include <vector>

using std::cout;
using std::generate;
using std::vector;

__global__ void matrixMul(const int *a, const int *b, int *c, int N) {
    // Compute each thread's global row and column index
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    // Iterate over row, and down column
    c[row * N + col] = 0;
    for (int k = 0; k < N; k++) {
        // Accumulate results for a single element
        c[row * N + col] += a[row * N + k] * b[k * N + col];
    }
}

// Check result on the CPU
void verify_result(vector<int> &a, vector<int> &b, vector<int> &c, int N) {
    for (int row = 0; row < N; row++) {
        for (int col = 0; col < N; col++) {
            int tmp = 0; // For every element in the row-column pair
            for (int k = 0; k < N; k++) {
                // Accumulate the partial results
                tmp += a[row * N + k] * b[k * N + col];
            }
            // Check against the CPU result
            assert(tmp == c[row * N + col]);
        }
    }
}

int main() {
    int N = 1 << 10; // Matrix size of 1024 x 1024;

    // Size (in bytes) of matrix
    size_t bytes = N * N * sizeof(int);

    // Host vectors
    vector<int> h_a(N * N);
    vector<int> h_b(N * N);
    vector<int> h_c(N * N);

    // Initialize matrices
    generate(h_a.begin(), h_a.end(), []() { return rand() % 100; });
    generate(h_b.begin(), h_b.end(), []() { return rand() % 100; });

    // Allocate device memory
    int *d_a, *d_b, *d_c;
    cudaMalloc(&d_a, bytes);
    cudaMalloc(&d_b, bytes);
    cudaMalloc(&d_c, bytes);

    // Copy data to the device
    cudaMemcpy(d_a, h_a.data(), bytes, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, h_b.data(), bytes, cudaMemcpyHostToDevice);

    // Threads per CTA dimension
    int THREADS = 32;

    // Blocks per grid dimension (assumes THREADS divides N evenly)
    int BLOCKS = N / THREADS;

    // Use dim3 structs for block and grid dimensions
    dim3 threads(THREADS, THREADS);
    dim3 blocks(BLOCKS, BLOCKS);

    // Launch kernel
    matrixMul<<<blocks, threads>>>(d_a, d_b, d_c, N);

    // Copy back to the host
    cudaMemcpy(h_c.data(), d_c, bytes, cudaMemcpyDeviceToHost);

    // Check result
    verify_result(h_a, h_b, h_c, N);

    cout << "COMPLETED SUCCESSFULLY\n";

    // Free memory on device
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    return 0;
}
```

Writing matrix_add.cu

```
In [2]: !echo "Check your GPU version"
!nvidia-smi

Check your GPU version
Sun Apr 16 17:09:24 2023

+-----+
| NVIDIA-SMI 510.108.03    Driver Version: 510.108.03    CUDA Version: 11.6    |
+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC | |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|              |                    |              | MIG M. |
+-----+-----+
| 0     NVIDIA GeForce ...   Off      | 00000000:01:00.0  On   |           N/A       | |
| 0%    56C    P5       29W / 250W   | 1465MiB / 8192MiB |      28%    Default  |
|              |                    |              | N/A     |
+-----+-----+

+-----+
| Processes: |
| GPU   GI   CI        PID   Type   Process name                      GPU Memory |
| ID   ID   ID              |                   |            Usage |
+-----+-----+
| 0     N/A   N/A       1703    G    /usr/lib/xorg/Xorg                  96MiB |
| 0     N/A   N/A       2624    G    /usr/lib/xorg/Xorg                  704MiB |
| 0     N/A   N/A       2819    G    /usr/bin/gnome-shell                91MiB |
| 0     N/A   N/A       3712    G    ...RendererForSitePerProcess        3MiB |
| 0     N/A   N/A       8175    G    ...features=BackForwardCache        10MiB |
| 0     N/A   N/A       8543    G    ...957248867340528764,131072       156MiB |
| 0     N/A   N/A      13016    G    ...AAAAAAAA= --shared-files         35MiB |
| 0     N/A   N/A      14339    G    ...b/thunderbird/thunderbird       137MiB |
| 0     N/A   N/A      15643    G    ...RendererForSitePerProcess       181MiB |
+-----+

+-----+
| Processes: |
| GPU   GI   CI        PID   Type   Process name                      GPU Memory |
| ID   ID   ID              |                   |            Usage |
+-----+-----+
| 0     N/A   N/A       1703    G    /usr/lib/xorg/Xorg                  96MiB |
| 0     N/A   N/A       2624    G    /usr/lib/xorg/Xorg                  704MiB |
| 0     N/A   N/A       2819    G    /usr/bin/gnome-shell                91MiB |
| 0     N/A   N/A       3712    G    ...RendererForSitePerProcess        3MiB |
| 0     N/A   N/A       8175    G    ...features=BackForwardCache        10MiB |
| 0     N/A   N/A       8543    G    ...957248867340528764,131072       156MiB |
| 0     N/A   N/A      13016    G    ...AAAAAAAA= --shared-files         35MiB |
| 0     N/A   N/A      14339    G    ...b/thunderbird/thunderbird       137MiB |
| 0     N/A   N/A      15643    G    ...RendererForSitePerProcess       181MiB |
+-----+
```

```
In [3]: %%bash

CUDA_SUFF=70 # or CUDA_SUFF=35
nvcc -gencode arch=compute_${CUDA_SUFF},code=sm_${CUDA_SUFF} ./matrix_add.cu -o matrix_add
./matrix_add
```

COMPLETED SUCCESSFULLY

```
In [4]: %%bash
# ls
# nvprof ./matrix_add
nvprof ./matrix_add

==28468== NVPROF is profiling process 28468, command: ./matrix_add
COMPLETED SUCCESSFULLY
==28468== Profiling application: ./matrix_add
==28468== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 79.46%  5.1380ms      1  5.1380ms  5.1380ms  5.1380ms  matrixMul(int const *,
int const *, int*, int)
              13.46%  870.09us      2  435.04us  429.16us  440.93us  [CUDA memcpy HtoD]
              7.08%  457.92us      1  457.92us  457.92us  457.92us  [CUDA memcpy DtoH]
              31.075ms  45.003us  93.136ms  cudaMalloc
              6.73%  6.7563ms      3  2.2521ms  470.65us  5.7748ms  cudaMemcpy
              0.32%  323.92us      3  107.97us  99.012us  113.98us  cudaFree
              0.07%  67.404us     101  667ns      84ns  28.975us  cuDeviceGetAttribute
              0.02%  21.152us      1  21.152us  21.152us  21.152us  cudaLaunchKernel
              0.02%  18.594us      1  18.594us  18.594us  18.594us  cuDeviceGetName
              0.00%  4.6960us      1  4.6960us  4.6960us  4.6960us  cuDeviceGetPCIBusId
              0.00%  1.3480us      3  449ns      125ns  1.0780us  cuDeviceGetCount
              0.00%  1.1780us      2  589ns      88ns  1.0900us  cuDeviceGet
              0.00%  330ns        1  330ns      330ns  330ns  cuDeviceTotalMem
              0.00%  257ns         1  257ns      257ns  257ns  cuDeviceGetUuid
```

What is the difference between ‘GPU activities’ and ‘API calls’ in the results of ‘nvprof’?

Answer from <https://forums.developer.nvidia.com/t/what-is-the-difference-between-gpu-activities-and-api-calls-in-the-results-of-nvprof/71338/1>

Section ‘GPU activities’ list activities which execute on the GPU like CUDA kernel, CUDA memcpy, CUDA memset. And timing information here represents the execution time on the GPU.

Section ‘API Calls’ list CUDA Runtime/Driver API calls. And timing information here represents the execution time on the host.

For example, CUDA kernel launches are asynchronous from the point of view of the CPU. It returns immediately, before the kernel has completed, and perhaps before the kernel has even started. This time is captured for the Launch API like cuLaunchKernel in the ‘API Calls’ section. Eventually kernel starts execution on the GPU and runs to the completion. This time is captured for kernel in the ‘GPU activities’.

```
In [5]: %%bash
nvprof --print-gpu-trace ./matrix_add --benchmark
```

```
==28485== NVPROF is profiling process 28485, command: ./matrix_add --benchmark
COMPLETED SUCCESSFULLY

==28485== Profiling application: ./matrix_add --benchmark
==28485== Profiling result:
   Start   Duration      Grid Size      Block Size  Regs*   SSMem*   DSMem*      Size  Thr
  oughput SrcMemType  DstMemType      Device  Context  Stream   Name      -      -
276.36ms  441.83us      -      -      -      -      -      -      -  4.0000MB  8.8
411GB/s   Pageable      Device  NVIDIA GeForce      1      7      [CUDA memcpy HtoD]
276.88ms  428.29us      -      -      -      -      -      -      -  4.0000MB  9.1
205GB/s   Pageable      Device  NVIDIA GeForce      1      7      [CUDA memcpy HtoD]
277.32ms  5.2855ms      (32 32 1)  (32 32 1)      28      0B      0B      -
-      -      -  NVIDIA GeForce      1      7      matrixMul(int const *, int con
st *, int*, int) [116]
282.60ms  444.36us      -      -      -      -      -      -      -  4.0000MB  8.7
908GB/s   Device      Pageable  NVIDIA GeForce      1      7      [CUDA memcpy DtoH]

Regs: Number of registers used per CUDA thread. This number includes registers used internally by th
e CUDA driver and/or tools and can be more than what the compiler shows.
SSMem: Static shared memory allocated per CUDA block.
DSMem: Dynamic shared memory allocated per CUDA block.
SrcMemType: The type of source memory accessed by memory operation/copy
DstMemType: The type of destination memory accessed by memory operation/copy
```

In []: